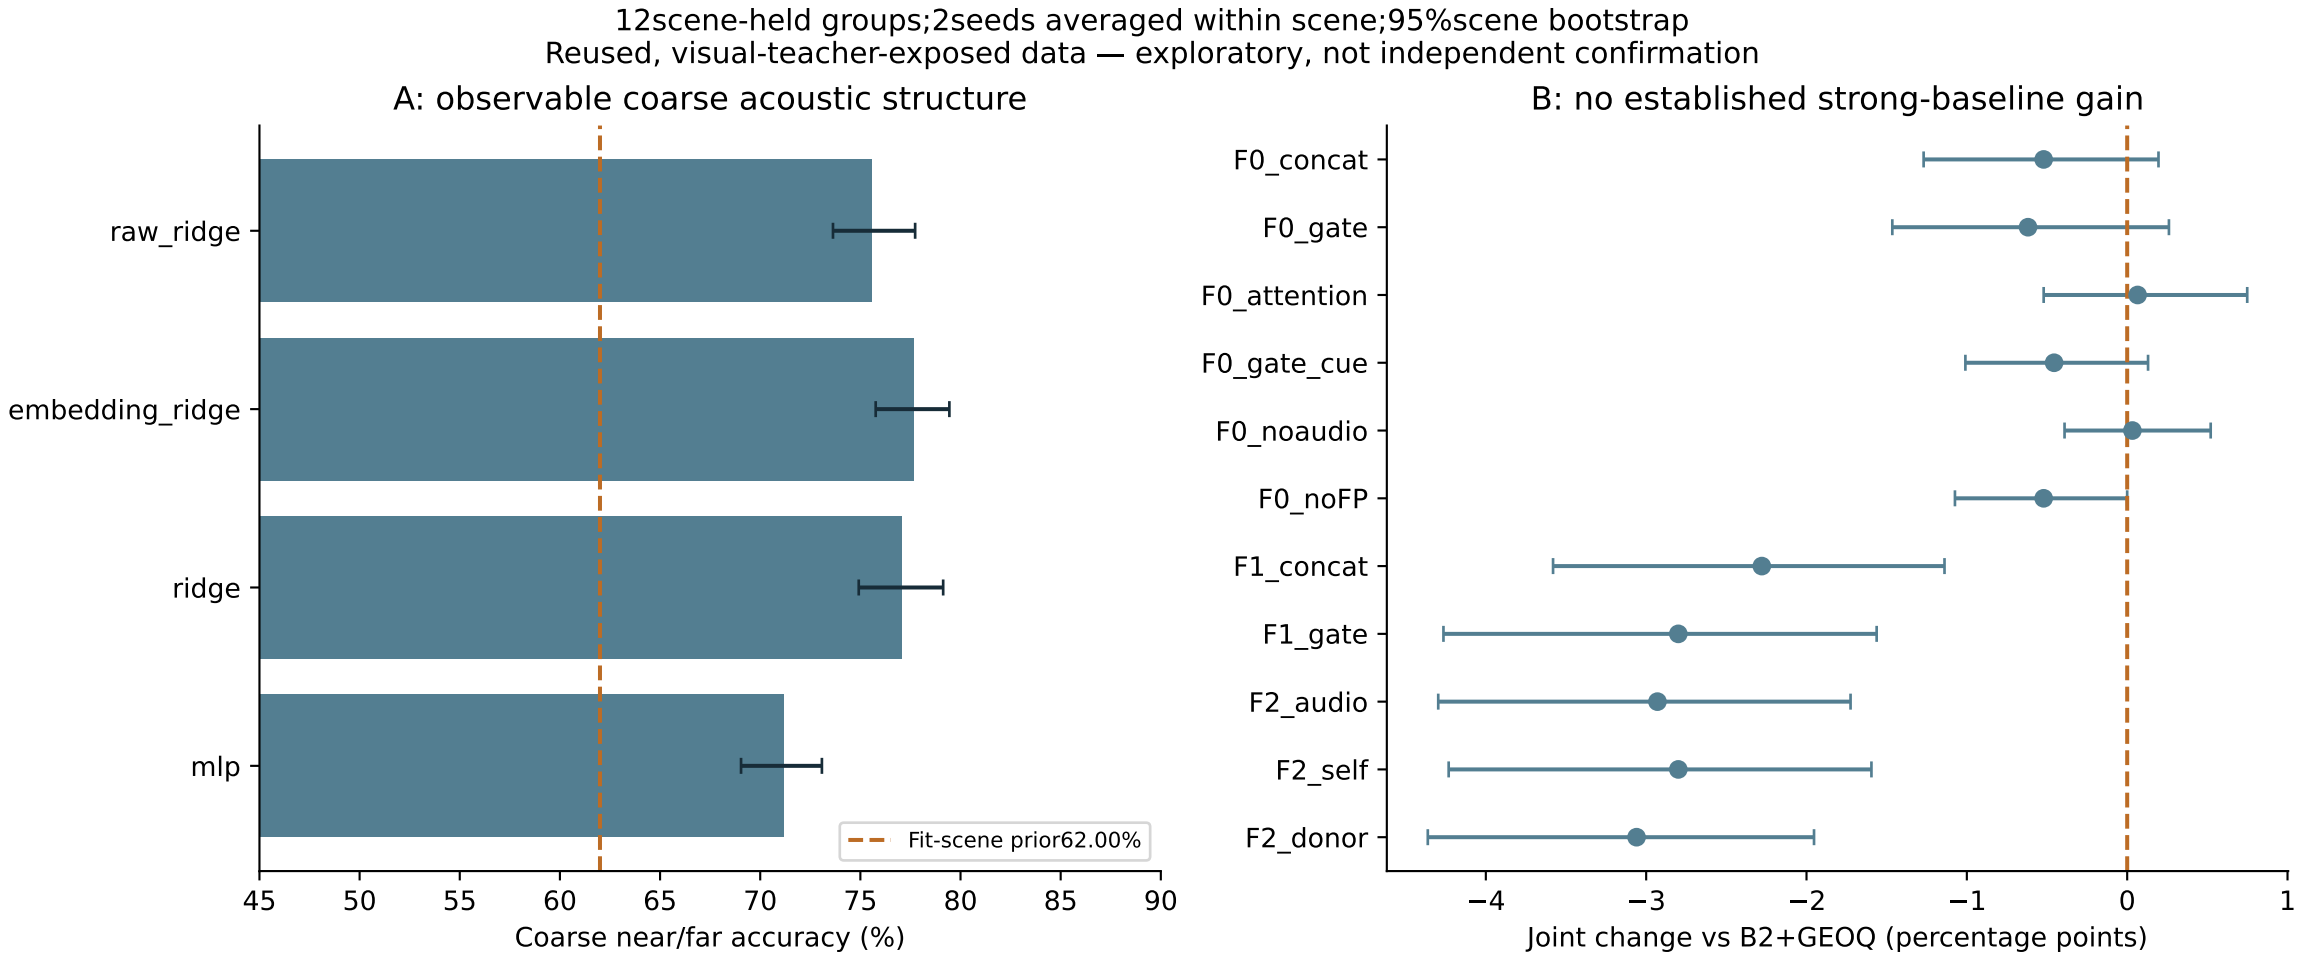
12scene-held groups;2seeds averaged within scene;95%scene bootstrap
Reused, visual-teacher-exposed data — exploratory, not independent confirmation
A: observable coarse acoustic structure
raw_ridge
embedding_ridge
ridge
mlp
Fit-scene prior62.00%
45
50
55
60
65
70
75
80
85
90
Coarse near/far accuracy (%)
B: no established strong-baseline gain
F0_concat
F0_gate
F0_attention
F0_gate_cue
F0_noaudio
F0_noFP
F1_concat
F1_gate
F2_audio
F2_self
F2_donor
−4
−3
−2
−1
0
1
Joint change vs B2+GEOQ (percentage points)